

WELCOME!!!

CS 6650 *Scalable* Distributed Systems
(while maintaining some sanity?!) 😊

What are we doing?

- Weekly class time
 - Yvonne talking fast but recorded!
 - Industry mentors visiting
 - Working together to get the *crazy code* to RUN on AWS (or where you want to?)!!!
- Weekly mock interviews (sign up!) to walk through your code with the teaching team!
- Please read the syllabus and ask any questions!

How are we going to do it?!

- Weekly homeworks and Masteries are reviewed in *mock interviews*
 - 1:1 time with your TA!
- Homework will cover
 - Basics of concurrency and distributed system architectures
 - Server processing
 - Data, Data and more Data
 - Scalability issues throughout
- We will spend class time on homeworks together!
- Canvas, Teams, ... talk to us!

The Homeworks build into larger Assignments!

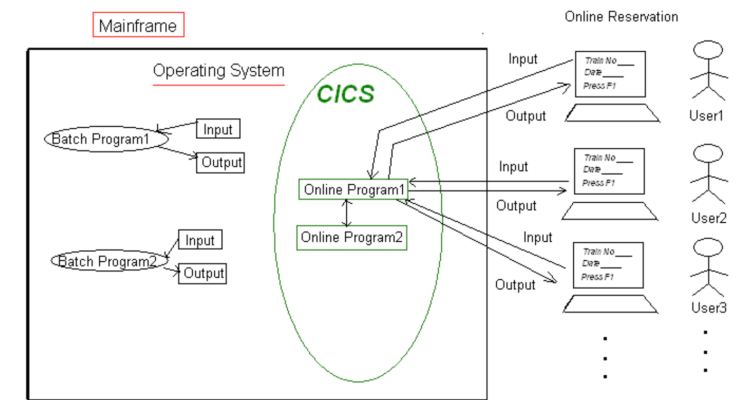
- Client
 - Multithreaded exerting request load on dumb server to start 😊
 - Infrastructure to measure is on client-side
- Server evolves!
 - Database introduced along with “business logic”
 - Data storage management
 - Write-oriented workloads (making it harder to scale)
 - Threading on server
 - State management, horizontal scaling

Server is Still evolving!

- Now scaling across more servers
 - Queuing for better responsiveness
 - Caching for performance
 - Monitoring
- And finally! 😊
 - Distributed database issues
 - More design patterns for data intensive computing
- Project ideas?!?

Quick intro, and history lesson!

- CICS IBM in the 60s! (still around)
 - Customer Information Control System
 - Mainframe and users
 - Timeshared server, dumb terminal
 - MTS (Michigan Terminal System)
- Along comes the Internet in the 90s
 - LANs, WANs, modems, and browsers (slow and simple!)
 - Eventually needed to start thinking about scale!
 - How many users/clients are hitting the banking server?
 - Amazon, Google, ebay?
 - Now we have MASSIVE data centers all over the world!
 - If we can process it!!!
 - Check out petabyte scale data sets
 - <https://www.lifewire.com/terabytes-gigabytes-amp-petabytes-how-big-are-they-4125169>



<http://www.mainframegurukul.com/tutorials/programming/cics/introduction-to-cics.html>

The screenshot shows a 'Make New Connection' dialog box. On the left is an illustration of a computer, a telephone, and a modem connected by lines. On the right, the text 'Type the phone number for the computer you want to call:' is followed by input fields. The 'Area code:' field contains '0844' and the 'Telephone number:' field contains 'xxxxxxx'. Below these is a 'Country code:' dropdown menu showing 'United Kingdom (44)'. At the bottom are three buttons: '< Back', 'Next >', and 'Cancel'.

Software fabric that make distributed systems!

- Remote Procedure Calls
 - Java RMI
- Sockets (80s!)
 - Websockets
- CORBA (remote objects)
- HTTP
- More evolution to come...
 - Languages like Go and Rust!
 - Dask?!
 - Projects! 😊

Applications today

- These are just a few...
- Numbers a little out of date
- BUT YOU GET THE IDEA!
- More applications coming
 - Your project?!
 - The Earth Data Store

Dropbox – [exascale storage system](#), millions of new files per hour, from 2012-16 seen over 12x growth.

Gmail – [1.2 billion users](#), blocks 10M spam messages every minute

Netflix – runs on AWS, [took 7 years to migrate](#), uses 15% of global downstream internet traffic

Facebook - [More than 300 million photos](#) uploaded per day and 510,000 comments posted and 293,000 statuses updated per minute

Youtube - Users watch [4,146,600 YouTube videos](#) every minute

What do these systems all have in common?

- Transactions, user interactions, analysis
- Many more have sensors, IoT
- Many more now have rich UIs
 - AR/VR!
- This means we have **large code** bases along with these massive data sets!
- And **large data sets** of performance monitoring results!
- They operate on multiple “nodes”, and are distributed systems

Nonfunctional Requirements (quality attributes)


- Performance, performance, performance
 - NOTE: collecting data on performance generates a lot of data!
 - Analysis of this is also challenging
- Scale
 - Performance is still good, even with more data, more users, more everything!
 - Note this can be very dynamic spikes
 - A Tweet from a well known person!
- Security
 - Data and access control
- Available
 - 24/7 (Leah knows about stress!)

What is SCALABILITY?

- The ability of a system to **increase** or **decrease** capacity in response to changing demands while continuing to satisfy service level agreements for latency and availability
- “Elastic Compute”
- Hyperscalability is the same as the above, but while supporting **exponential growth** with **linear costs**

Can we just throw hardware at it?

- Yes and no!
- We want linear “scale up”
 - With each additional machine, we want to be able to grow linearly
 - BUT there is Amdahl’s Law (wikipedia here!)
 - At some point we level out
 - Not getting faster because some things have to be done in sequence, one at a time
 - Example, a write to a database
 - “Serial bottlenecks”!




SOME problems are easily parallelizable

- Map/Reduce (Google paper, please share your thoughts on Piazza!)
- Some serialized behaviour but does well!
- Can be a little tricky to get right...
 - Data Skew
 - Breaking the data set down into “chunks”
 - Not every chunk takes the same amount of time to execute
- Please take a look at the paper (link provided on Canvas) and share your thoughts!

Response times are critical!

- Within *40 milliseconds*
 - A user will LEAVE!
 - Impacts google rankings!
 - There are delays in data transfer (latencies per request)
 - You will love measuring these! 😊
- Fun fact is that there is a VERY LONG TAIL graph
 - Slow responses DO happen
 - So how can we mitigate those?!



Throughput

- Typically reported as *requests per second*
 - Note: Latency is for a single request
 - Throughput is a bulk measure of **many** requests
 - You will love measuring this!
- When you have millions of requests per second *and* bursts
 - We need to split the load across multiple nodes!
- “Scale-out”
 - Add more nodes to increase throughput...
 - Does it always work?

Availability!

- 24/7
 - 99%, 99.999%, 99.9999%...

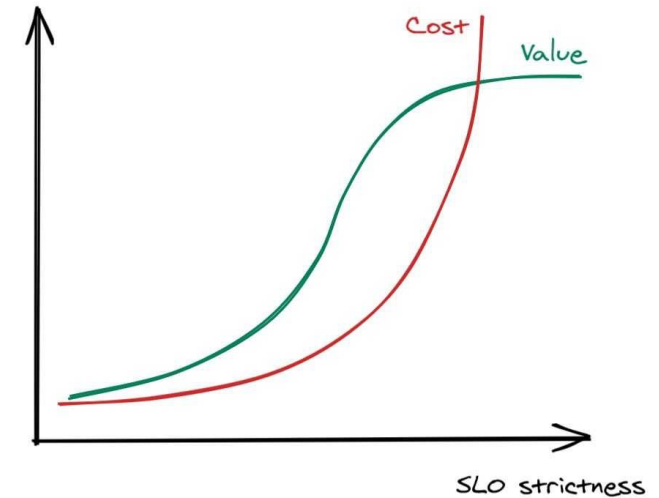
Availability %	Downtime per year ^[note 1]	Downtime per quarter	Downtime per month	Downtime per week	Downtime per day (24 hours)
90% ("one nine")	36.53 days	9.13 days	73.05 hours	16.80 hours	2.40 hours
95% ("one and a half nines")	18.26 days	4.56 days	36.53 hours	8.40 hours	1.20 hours
97%	10.96 days	2.74 days	21.92 hours	5.04 hours	43.20 minutes
98%	7.31 days	43.86 hours	14.61 hours	3.36 hours	28.80 minutes
99% ("two nines")	3.65 days	21.9 hours	7.31 hours	1.68 hours	14.40 minutes
99.5% ("two and a half nines")	1.83 days	10.98 hours	3.65 hours	50.40 minutes	7.20 minutes
99.8%	17.53 hours	4.38 hours	87.66 minutes	20.16 minutes	2.88 minutes
99.9% ("three nines")	8.77 hours	2.19 hours	43.83 minutes	10.08 minutes	1.44 minutes
99.95% ("three and a half nines")	4.38 hours	65.7 minutes	21.92 minutes	5.04 minutes	43.20 seconds
99.99% ("four nines")	52.60 minutes	13.15 minutes	4.38 minutes	1.01 minutes	8.64 seconds
99.995% ("four and a half nines")	26.30 minutes	6.57 minutes	2.19 minutes	30.24 seconds	4.32 seconds
99.999% ("five nines")	5.26 minutes	1.31 minutes	26.30 seconds	6.05 seconds	864.00 milliseconds
99.9999% ("six nines")	31.56 seconds	7.89 seconds	2.63 seconds	604.80 milliseconds	86.40 milliseconds
99.99999% ("seven nines")	3.16 seconds	0.79 seconds	262.98 milliseconds	60.48 milliseconds	8.64 milliseconds
99.999999% ("eight nines")	315.58 milliseconds	78.89 milliseconds	26.30 milliseconds	6.05 milliseconds	864.00 microseconds
99.9999999% ("nine nines")	31.56 milliseconds	7.89 milliseconds	2.63 milliseconds	604.80 microseconds	86.40 microseconds

https://en.wikipedia.org/wiki/High_availability

- Don't want a single point of failure (SPoF)
 - So we need to make sure we have replicas/redundancy!
- When you have 1000s of components, some will fail!
 - Need to incorporate FAILURE into our designs

It is all about the *Tradeoffs*!

- Throughput versus cost
 - More hardware means more \$\$
 - When is it worth it? When isn't it?
- Performance versus availability
 - Fast but components may fail!
 - Remember the old mainframes, they were fast-ish 😊
- ***Need to be able to measure and discuss the tradeoffs in the work you do in this course!***
- This all depends on the context of the requirements of the system you are building, there is no ONE right answer!



<https://robertovitillo.com/why-you-should-measure-tail-latencies/>